

On Gradient

Giulio Galvan

17 marzo 2015

Sommario

1 How to compute gradient

1.1 Backpropagation

First of all we need to define a loss function over the training data, so we define a dataset as

$$D = \{x^{(i)} \in \mathbb{R}^p, y^{(i)} \in \mathbb{R}^q, i \in [1, N]\} \quad (1)$$

and the loss function as

$$L(W) = \frac{1}{N} \sum_{i=1}^N L^{(i)}(W) \quad (2)$$

where W represents all the weights of the net. The network is defined by

$$a_l \triangleq \sum_j w_{lj} \phi_j \quad (3)$$

$$\phi_l \triangleq \sigma(a_l) \quad (4)$$

where w_{lj} is the weight of the connection between neuron j and neuron l and σ is the non linear activation function

So we can compute partial derivatives with respect to a single weight w_{lj} , using simply the chain rule, as

$$\frac{\partial L^{(i)}}{\partial w_{lj}} = \frac{\partial L^{(i)}}{\partial a_l} \cdot \frac{\partial a_l}{\partial w_{lj}} = \delta_l \cdot \phi_j$$

where we put

$$\delta_l \triangleq \frac{\partial L^{(i)}}{\partial a_l} \quad (5)$$




Figura 1: Modello per RNN

So we can easily compute $\delta_u = \frac{\partial L^{(i)}}{\partial a_u}$ for each output unit u once we choose a differentiable loss function; note that we don't need the weights for such a computation.

Let $P(l)$ be the set of parents of neuron l , formally:

$$P(l) = \{k : \exists \text{ a link between } l \text{ and } k \text{ with weight } w_{lk}\} \quad (6)$$

Again, simply using the chain rule, we can write, for each non output unit l :

$$\delta_l = \sum_{k \in P(l)} \frac{\partial L^{(i)}}{\partial a_k} \cdot \frac{\partial a_k}{\partial a_l} = \sum_{k \in P(l)} \delta_k \cdot \frac{\partial a_k}{\partial \phi_l} \cdot \frac{\partial \phi_l}{\partial a_l} = \sum_{k \in P(l)} \delta_k \cdot w_{kl} \cdot \sigma'(a_l) \quad (7)$$

For output units instead we can compute $\delta_u = \frac{\partial L^{(i)}}{\partial a_u}$ directly once we define the loss function.

1.2 Backpropagation matrix notation

Here we rewrite the previously derived equations in matrix notation.

Let us define the weight matrix $W_i \in \mathbb{R}_{(p(i), p(i-1))}$, whose element $W_{i,j}$ is the weight of the arc which links neuron j from level $i-1$ to neuron i from level i , where $p(i)$ is the neuron number for i^{th} level.

$$\vec{\phi}_1 \triangleq \vec{x} \quad (8)$$

$$\vec{a}_{i+1} \triangleq W_{i+1} \cdot \vec{\phi}_i \quad (9)$$

$$\vec{\phi}_{i+1} \triangleq \sigma(\vec{a}_{i+1}) \quad (10)$$

where $\sigma(\cdot)$ is the non-linear activation function and it's applied element by element. We can rewrite equation 7 in matrix notation as:

$$\frac{\partial L}{\partial W_i} = \frac{\partial L}{\partial \vec{a}_i} \cdot \frac{\partial \vec{a}_i^T}{\partial W_i} = \Delta_i \cdot \vec{\phi}_{i-1}^T \quad (11)$$

where

$$\Delta_i \triangleq \frac{\partial L}{\partial \vec{a}_i} \quad (12)$$

$$\Delta_i = W_{i+1}^T \cdot \Delta_{i+1} \circ \sigma(\Delta_i) \quad (13)$$

Riferimenti bibliografici

- [1] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory, 1995.